

1NT19IS108
CHAITANYA P
C1 BATCH

```
hadoop@admin1-HP-280-G4-MT-Business-PC: ~/apache-hive-3.1.2-bin/conf
hadoop@admin1-HP-280-G4-MT-Business-PC:~$ cd $HADOOP_HOME
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1$ cd sbin/
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ ./start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [admin1-HP-280-G4-MT-Business-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ jps
3864 DataNode
4296 ResourceManager
4824 Jps
3723 NameNode
4108 SecondaryNameNode
4447 NodeManager
hadoop@admin1-HP-280-G4-MT-Business-PC:~/hadoop-3.2.1/sbin$ cd $HIVE_HOME
hadoop@admin1-HP-280-G4-MT-Business-PC:~/apache-hive-3.1.2-bin$ ./conf
bash: ./conf: Is a directory
hadoop@admin1-HP-280-G4-MT-Business-PC:~/apache-hive-3.1.2-bin$ cd conf
hadoop@admin1-HP-280-G4-MT-Business-PC:~/apache-hive-3.1.2-bin/conf$ ls
beeline-log4j2.properties.template
derby.log
hive-default.xml.template
hive-env.sh.template
hive-exec-log4j2.properties.template
hive-log4j2.properties.template
ivysettings.xml
llap-cli-log4j2.properties.template
llap-daemon-log4j2.properties.template
metastore_db
parquet-logging.properties
hadoop@admin1-HP-280-G4-MT-Business-PC:~/apache-hive-3.1.2-bin/conf$ rm -rf metastore_db
hadoop@admin1-HP-280-G4-MT-Business-PC:~/apache-hive-3.1.2-bin/conf$ schematool -initSchema -dbType derby
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
```

Dropping database
Rm -rf metastore_db

```
parquet-logging.properties
hadoop@admin1-HP-280-G4-MT-Business-PC:~/apache-hive-3.1.2-bin/conf$ rm -rf metastore_db
hadoop@admin1-HP-280-G4-MT-Business-PC:~/apache-hive-3.1.2-bin/conf$ schematool -initSchema -dbType derby
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
```

```
hadoop@admin1-HP-280-G4-MT-Business-PC: ~/apache-hive-3.1.2-bin/conf
hadoop@admin1-HP-280-G4-MT-Business-PC:~/apache-hive-3.1.2-bin/conf$ hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/hadoop/hadoop-3.2.1/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = a36d5c8e-d0d8-41a2-9832-9af9457d4436

Logging initialized using configuration in jar:file:/home/hadoop/apache-hive-3.1.2-bin/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive Session ID = 3e60ae10-704c-470e-acf4-2b988cbcd39a
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
hive> show databases;
OK
default
Time taken: 0.446 seconds, Fetched: 1 row(s)
hive> create database if not exists salesdb;
OK
Time taken: 0.076 seconds
hive> use salesdb;
OK
Time taken: 0.03 seconds
hive> create table sales(product string,price int,payment_type string,name string,city string,state string,country string);
OK
Time taken: 0.502 seconds
hive> load data local inpath '/home/hadoop/sales.csv' into table sales;
FAILED: SemanticException Line 1:23 Invalid path ''/home/hadoop/sales.csv': No files matching path file:/home/hadoop/sales.csv
hive> load data local inpath '/home/hadoop/Desktop/sales.csv' into table sales;
Loading data to table salesdb.sales
OK
Time taken: 0.859 seconds
hive> select * from sales;
OK
,,,,, NULL NULL NULL NULL NULL NULL
Product1,1200,Mastercard,carolina,Basildon,England,United Kingdom NULL NULL NULL NULL NULL NULL
Product1,1200,Visa,Betina,Parkville,MO,United States NULL NULL NULL NULL NULL NULL
```

show databases;

create database if not exists chaitanyadb;

use chaitanyadb;

```
hive> show databases;
OK
default
salesdb
Time taken: 0.039 seconds, Fetched: 2 row(s)
hive> create database if not exists chaitanyadb;
OK
Time taken: 0.102 seconds
hive> use chaitanyadb;
OK
Time taken: 0.033 seconds
hive> create table sales(product string,price int,payment_type string,name string,city string,state string,country string);
OK
Time taken: 0.094 seconds
hive>
> insert into sales values('product1,5000','Mastercard','harry','Astoria','California','USA');
MismatchedTokenException(241=374)
    at org.apache.hadoop.hive.ql.parse.HiveParser.recoverFromMismatchedToken(BaseRecognizer.java:617)
    at org.apache.hadoop.hive.ql.parse.HiveParser.match(BaseRecognizer.java:115)
    at org.apache.hadoop.hive.ql.parse.HiveParser.expressionsInParenthesis(HiveParser_IdentifierParser.java:2316)
    at org.apache.hadoop.hive.ql.parse.HiveParser.expressionsInParenthesis(HiveParser.java:45260)
    at org.apache.hadoop.hive.ql.parse.HiveParser.FromClauseParser.valueRowConstructor(HiveParser_FromClauseParser.java:6214)
    at org.apache.hadoop.hive.ql.parse.HiveParser.FromClauseParser.valuesTableConstructor(HiveParser_FromClauseParser.java:6131)
```

Creating table:

create table sales(product string,price int,payment_type string,name string,city string,state string,country string);

insert into sales values('product1,5000','Mastercard','harry','Astoria','California','USA');

insert into sales values('product2',1000,'Visa','ron','Banglore','Karnataka','India');

```

insert into sales values('product3',3000,'Diners','draco','Banglore','Karnataka','India');
insert into sales values('product4',3000,'MasterCard','luna','Banglore','Karnataka','India');
insert into sales values('product1',5000,'Visa','dove','Chittor','AP','India');
insert into sales values('product2',6000,'Rupay','elena','ATP','AP','India');

```

```

hive> select * from sales;
OK
product1      5000    Mastercard    harry    Astoria    California    USA
product2      1000    Visa         ron      Banglore    Karnataka     India
product3      3000    Diners       draco     Banglore    Karnataka     India
product4      3000    MasterCard    luna     Banglore    Karnataka     India
product1      5000    Visa         dove     Chittor    AP            India
product2      6000    Rupay        elena     ATP         AP            India
Time taken: 0.07 seconds, Fetched: 6 row(s)

```

```

hadoop@admin1-HP-280-G4-MT-Business-PC: ~/apache-hive-3.1.2-bin/conf
OK
Time taken: 16.996 seconds
hive> insert into sales values('product4',3000,'MasterCard','luna','Banglore','Karnataka','India');
Query ID = hdoop_20220628115145_7c4a1ca7-3904-4d92-9ce2-70f9a787e471
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1656395330444_0008, Tracking URL = http://admin1-HP-280-G4-MT-Business-PC:8088/proxy/application_1656395330444_0008/
Kill Command = /home/hadoop/hadoop-3.2.1/bin/mapred job -kill job_1656395330444_0008
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-28 11:51:52,055 Stage-1 map = 0%, reduce = 0%
2022-06-28 11:51:57,235 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.54 sec
2022-06-28 11:52:01,358 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.07 sec
MapReduce Total cumulative CPU time: 4 seconds 70 msec
Ended Job = job_1656395330444_0008
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to directory hdfs://127.0.0.1:9000/user/hive/warehouse/chaitanyadb.sales/.hive-staging_hive_2022-06-28_11-51-45_523_4289236211524
882476-1/-ext-10000
Loading data to table chaitanyadb.sales
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.07 sec HDFS Read: 21684 HDFS Write: 470 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 70 msec
OK
Time taken: 18.217 seconds
hive> insert into sales values('product1',5000,'Visa','dove','Chittor','AP','India');
Query ID = hdoop_20220628115206_5fcf2c56-cfeb-49a0-98b4-21ec50dc4838
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>

```

Dropping the table:

```

create table Salestb(Product string,Price int,Payment_type string,Name string,City string,State
string,Country string)
drop table Salestb;

```

```
hive> create table Salestb(Product string,Price int,Payment_type string,Name string,City string,State string,Country string);
OK
Time taken: 0.056 seconds
hive> drop table Salestb;
OK
Time taken: 0.841 seconds
hive> █
```

Alter table:

alter table sales rename to chaitanya_sales;

```
hive> alter table sales rename to chaitanya_sales;
OK
Time taken: 0.137 seconds
hive> show tables;
OK
chaitanya_sales
Time taken: 0.014 seconds, Fetched: 1 row(s)
hive> █
```

3.Built in operators:

Arithmetic Operator in Hive

select product,name,price+1100 from sales;

```
chaitanya_sales
Time taken: 0.014 seconds, Fetched: 1 row(s)
hive> select product,name,price+1100 from chaitanya_sales;
OK
product1      harry      6100
product2      ron        2100
product3      draco      4100
product4      luna       4100
product1      dove       6100
product2      elena      7100
Time taken: 0.097 seconds, Fetched: 6 row(s)
hive> █
```

```
hive> select product,name,price+1100 from sales;
OK
product1      harry      6100
product2      ron        2100
product3      draco      4100
product4      luna       4100
product1      dove       6100
product2      elena      7100
Time taken: 0.187 seconds, Fetched: 6 row(s)
hive> █
```

select product,name,(price*10)/2 from sales;

```
hive> select product,name,(price*10)/2 from sales;
OK
product1      harry      25000.0
product2      ron        5000.0
product3      draco      15000.0
product4      luna       15000.0
product1      dove       25000.0
product2      elena      30000.0
Time taken: 0.1 seconds, Fetched: 6 row(s)
```

Relational Operators in Hive

select * from sales where price>4000;

```
hive> select * from int19is107_sales where price>4000;
FAILED: SemanticException [Error 10001]: Line 1:14 Table not found 'int19is107_sales'
hive> select * from sales where price>4000;
OK
product1      5000      mastercard      harry      astoria      california      USA
product1      5000      Visa      dove      Chittor      AP      India
product2      6000      Rupay      elena      ATP      AP      India
Time taken: 0.099 seconds, Fetched: 3 row(s)
```

4.Built in functions:

Mathematical Functions in Hive

select product,name,sqrt(price) from sales;

```
hive> select product,name,sqrt(price) from sales;
OK
product1      harry      70.71067811865476
product2      ron        31.622776601683793
product3      draco      54.772255750516614
product4      luna       54.772255750516614
product1      dove       70.71067811865476
product2      elena      77.45966692414834
Time taken: 0.103 seconds, Fetched: 6 row(s)
```

select product,name,floor(price) from sales;

```

Time taken: 0.103 seconds, Fetched: 6 row(s)
hive> select product,name,floor(price) from sales;
OK
product1      harry      5000
product2      ron        1000
product3      draco      3000
product4      luna       3000
product1      dove       5000
product2      elena      6000
Time taken: 0.107 seconds, Fetched: 6 row(s)
hive> select product,name,cos(price) from sales;
OK

```

select product,name,cos(price) from sales;

```

Time taken: 0.107 seconds, Fetched: 6 row(s)
hive> select product,name,cos(price) from sales;
OK
product1      harry      0.15466840618074712
product2      ron        0.562379076290703
product3      draco      -0.9756821998857504
product4      luna       -0.9756821998857504
product1      dove       0.15466840618074712
product2      elena      0.9039115103477952
Time taken: 0.117 seconds, Fetched: 6 row(s)
hive> select max(price) from sales where state=Karn

```

Aggregate Functions in Hive

select max(price) from sales where state=Karnataka;

```

hive> select max(price) from sales where state='Karnataka';
Query ID = hdoop_20220628134704_026eb5a1-00fc-477f-b8ef-56377f113a08
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1656401979434_0007, Tracking URL = http://admin1-HP-280-G4-MT-Business-PC:8088/proxy/application_1656401979434_0007/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1656401979434_0007
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-28 13:47:10,006 Stage-1 map = 0%, reduce = 0%
2022-06-28 13:47:14,116 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.3 sec
2022-06-28 13:47:19,232 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.89 sec
MapReduce Total cumulative CPU time: 3 seconds 890 msec
Ended Job = job_1656401979434_0007
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.89 sec HDFS Read: 15461 HDFS Write: 104 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 890 msec

```

Other built-in functions

select product, upper(name) from sales;

```
hive> select product, upper(name) from sales;
OK
product1      HARRY
product2      RON
product3      DRACO
product4      LUNA
product1      DOVE
product2      ELENA
Time taken: 0.075 seconds, Fetched: 6 row(s)
```

select product, reverse(name) from sales;

```
Time taken: 0.075 seconds, Fetched: 6 row(s)
hive> select product, reverse(name) from sales;
OK
product1      yrrah
product2      nor
product3      ocard
product4      anul
product1      evod
product2      anele
Time taken: 0.098 seconds, Fetched: 6 row(s)
```

select product, concat(name,state) from sales;

```
Time taken: 0.098 seconds, Fetched: 6 row(s)
hive> select product, concat(name,state) from sales;
OK
product1      harrycalifornia
product2      ronKarnataka
product3      dracoKarnataka
product4      lunaKarnataka
product1      doveAP
product2      elenaAP
```

select product, trim(name) from sales;

```

Time taken: 0.096 seconds, Fetched: 6 row(s)
hive> select product, trim(name) from sales;
OK
product1      harry
product2      ron
product3      draco
product4      luna
product1      dove
product2      elena
Time taken: 0.089 seconds, Fetched: 6 row(s)
hive> create view visa as select * from sales

```

5.VIEWS:

create view visa as select * from sales where payment_type='Visa';
select * from visa;

```

hive> create view visa as select * from sales where payment_type='Visa';
OK
Time taken: 0.1 seconds
hive> select * from visa;
OK
product2      1000      Visa      ron      Bangalore      Karnataka      India
product1      5000      Visa      dove      Chittor AP      India
Time taken: 0.1 seconds, Fetched: 2 row(s)
hive>

```

create view mastercard as select * from sales where payment_type=MasterCard;
select * from mastercar;

```

hive> create view mastercar as select * from sales where payment_type='MasterCard';
OK
Time taken: 0.124 seconds
hive> select * from mastercar;
OK
product4      3000      MasterCard      luna      Bangalore      Karnataka      India
Time taken: 0.102 seconds, Fetched: 1 row(s)

```

6.HIVEQL

GROUPBY:

select count(*)as count,country from sales group by country;


```

hive> select count(*) as count, country from sales group by country;
Query ID = hdoop_20220628140240_56837ef9-a2da-4c70-b264-587855bba5de
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1656401979434_0008, Tracking URL = http://admin1-HP-280-G4-MT-Business-PC:8088/proxy/application_1656401979434_0008/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1656401979434_0008
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-28 14:02:45,939 Stage-1 map = 0%, reduce = 0%
2022-06-28 14:02:50,045 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.83 sec
2022-06-28 14:02:55,169 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.04 sec
MapReduce Total cumulative CPU time: 3 seconds 40 msec
Ended Job = job_1656401979434_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.04 sec HDFS Read: 14708 HDFS Write: 125 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 40 msec
OK
5      India
1      USA
Time taken: 16.082 seconds, Fetched: 2 row(s)

```

select count(*) as count, state from sales group by state;

```

hive> select count(*) as count, state from sales group by state;
Query ID = hdoop_20220628140302_df4ddb8-df09-4de8-8ee2-9efad3ed66b8
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1656401979434_0009, Tracking URL = http://admin1-HP-280-G4-MT-Business-PC:8088/proxy/application_1656401979434_0009/
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1656401979434_0009
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-28 14:03:08,797 Stage-1 map = 0%, reduce = 0%
2022-06-28 14:03:12,973 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.83 sec
2022-06-28 14:03:17,060 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.29 sec
MapReduce Total cumulative CPU time: 3 seconds 290 msec
Ended Job = job_1656401979434_0009
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.29 sec HDFS Read: 14800 HDFS Write: 153 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 290 msec
OK

```

```

Total MapReduce CPU Time Spent: 3 seconds 290 msec
OK
2      AP
3      Karnataka
1      california
Time taken: 16.216 seconds, Fetched: 3 row(s)

```

select product, name from sales group by product, name;

```

hive> select product,name from sales group by product,name;
Query ID = hdoop_20220628140325_b651616c-8164-4b3c-a616-dbec8ef8bd3b
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1656401979434_0010, Tracking URL = http://admin1-HP-280-G4-MT-Business-PC:8088/proxy,
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1656401979434_0010
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-28 14:03:31,074 Stage-1 map = 0%, reduce = 0%
2022-06-28 14:03:35,226 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.55 sec
2022-06-28 14:03:40,378 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.02 sec
MapReduce Total cumulative CPU time: 3 seconds 20 msec
Ended Job = job_1656401979434_0010
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.02 sec HDFS Read: 13723 HDFS Write: 245 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 20 msec
OK
product1      dove
product1      harry
product2      elena
product2      ron
product3      draco
product4      luna

```

WHERE:

select * from sales where state=Karnataka;

```

hive> select * from sales where state='Karnataka';
OK
product2      1000    Visa    ron    Bangalore    Karnataka    India
product3      3000    Diners  draco    Bangalore    Karnataka    India
product4      3000    MasterCard  luna    Bangalore    Karnataka    India
Time taken: 0.063 seconds, Fetched: 3 row(s)
hive> select * from sales order by price;

```

ORDERBY:

select * from sales order by price;

```

hive> select * from sales order by price;
Query ID = hdoop_20220628140413_0b0554a2-7ff2-4621-b870-681e45d24e4d
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1656401979434_0011, Tracking URL = http://admin1-HP-280-G4-MT-Business-PC:8088/proxy/application_165
Kill Command = /home/hdoop/hadoop-3.2.1/bin/mapred job -kill job_1656401979434_0011
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2022-06-28 14:04:19,375 Stage-1 map = 0%, reduce = 0%
2022-06-28 14:04:23,492 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.41 sec
2022-06-28 14:04:28,629 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.75 sec
MapReduce Total cumulative CPU time: 2 seconds 750 msec
Ended Job = job_1656401979434_0011
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.75 sec HDFS Read: 13315 HDFS Write: 448 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 750 msec
OK
product2      1000    Visa    ron    Bangalore    Karnataka    India
product4      3000    MasterCard    luna    Bangalore    Karnataka    India
product3      3000    Diners    draco    Bangalore    Karnataka    India
product1      5000    Visa    dove    Chittor AP    India
product1      5000    mastercard    harry    astoria    california    USA
product2      6000    Rupay    elena    ATP    AP    India
Time taken: 15.922 seconds, Fetched: 6 row(s)
hive>

```